



# Logistics

- Lab 1 – data preparation
  - Instructions posted
  - In-class Lab Friday (XML, JSON, HTML)
  - Homework Due Friday 11:59pm
- Piazza Signup

# Review

- Data Types and Sources
  - Tabular
  - Semi-structured
  - Text, video, images
  - Graph data
- Data Models
- Data Preparation

# Data Model and Schema

A *data model* is a collection of concepts for describing data.


A *schema* is a description of a particular collection of data, using a given data model.

# Some common data models

- The **relational** model of data is the most widely used for record keeping.
  - Main concept: relation, basically a table with rows and columns.
  - Every relation has a schema, which describes the columns
- Semi-structured models in increasing use (e.g. XML)
  - Main concept: self-describing (tagged) document, basically a textual hierarchy (tree) of labeled values
  - Document Type Definition (DTD) or Schema possible, but not required
- Free text (and hypertext) widely used as well
  - Data represented for human consumption
    - Visual aspects and linguistic subtlety more important than clearly structured data
- Others: RDF Triple, Graph, Streaming, Probabilistic Data, Key-Value, Array/Matrix, Column Stores

# The Relational Model\*

- A Data Model based on Set/Bag Theory
  - Support Relational Algebra
- The Relational Model is Ubiquitous:
  - MySQL, PostgreSQL, Oracle, DB2, SQLServer, ...
  - Foundational work done at
    - IBM - System R
    - UC Berkeley - Ingres
- Object-oriented concepts have been merged in
  - Early work: POSTGRES research project at Berkeley
  - Informix, IBM DB2, Oracle 8i



E. F., “Ted” Codd  
Turing Award 1981

*\*Codd, E. F. (1970). "A relational model of data for large shared data banks".  
Communications of the ACM 13 (6): 37*

# Relational Database: Definitions

- *Relational database*: a set of *relations*
- *Relation*: made up of 2 parts:

*Schema* : specifies name of relation, plus name and type of each column

**Students**(*sid*: string, *name*: string, *login*: string, *age*: integer, *gpa*: real)

*Instance* : the actual data at a given time

- #rows = *cardinality*
- #fields = *degree / arity*

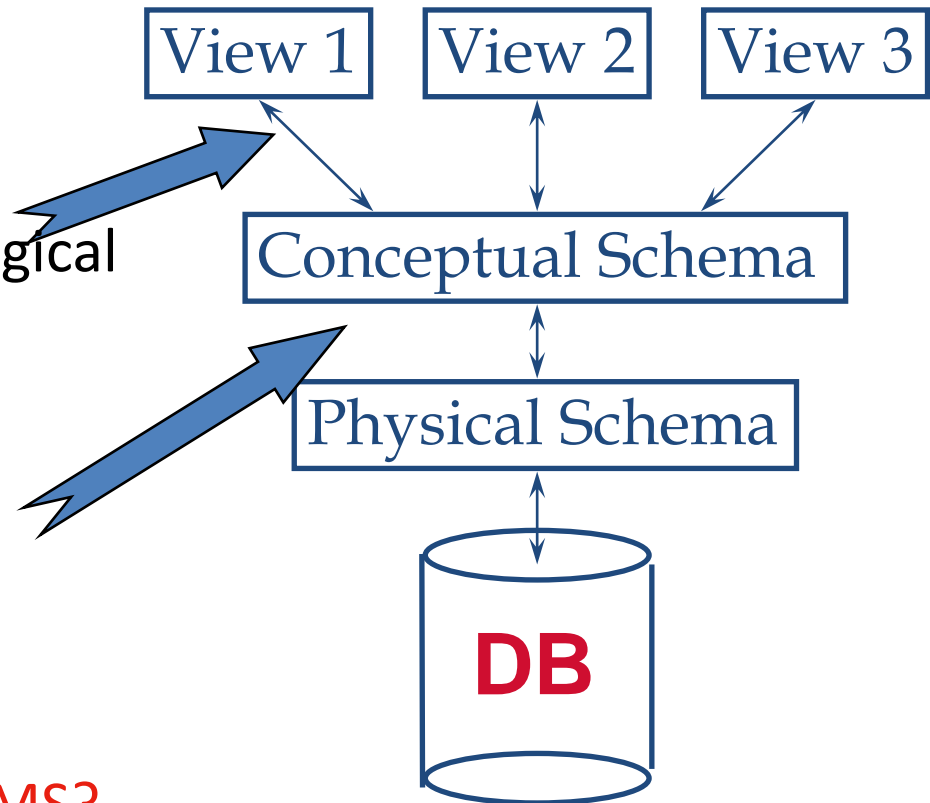
# Ex: Instance of Students Relation

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith @math	19	3.8

- Cardinality = 3, arity = 5 , all rows distinct
- The relation is true for these tuples and false for others (a.k.a, the closed world assumption)

# Levels of Abstraction

- Applications insulated from how data is structured and stored.
- **Logical data independence:**  
Protection from changes in logical structure of data.
- **Physical data independence:**  
Protection from changes in physical structure of data.
- **Q: Why are these particularly important for DBMS?**  
Because rate of change of DB applications is incredibly slow.





# Other Table-like Data Models:

## Pandas/Python

- **Series:** a named, ordered array/dictionary
  - **Values** can be any Numpy data type object
  - The **keys** of the dictionary are the **indexes**
  - Built on NumPy's **ndarray**
- **DataFrame:** a table with named columns
  - Represented as a **map** Dict (col\_name -> series)
  - Each Series object represents a **column**
- SQL vs. Pandas / SQL + Pandas

# Cases where Tables break as a data model? E.g., semi-/un-structured data



- Too limited in structure?
- Too rigid?
- Too costly up front?
- Examples... ??

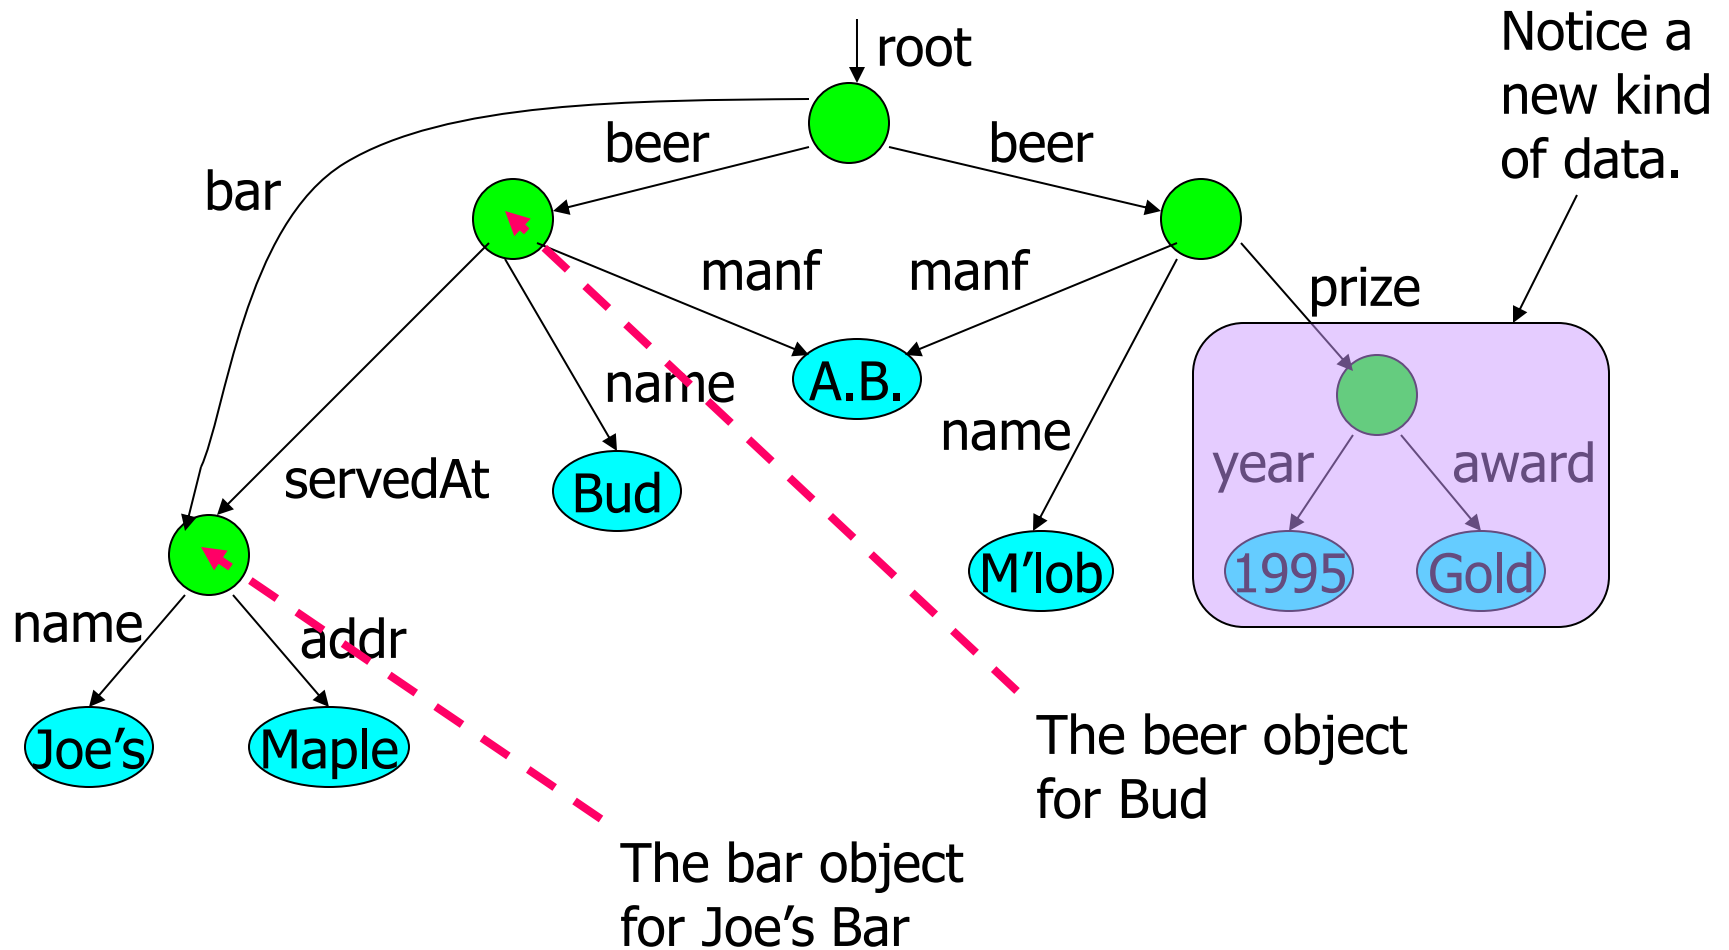
# Semistructured Data

- Another data model, based on trees.
- **Motivation:** flexible representation of data.
  - Often, data comes from multiple sources with differences in notation, meaning, etc.
- **Motivation:** sharing of *documents* among systems and databases.

# Hierarchical Semistructured Data

- Nodes = objects.
- Labels on arcs (attributes, relationships).
- Atomic values at leaf nodes (nodes with no arcs out).
- Flexibility: no restriction on:
  - Labels out of a node.
  - Number of successors with a given label.

# Example: Data Graph



# Well-Formed and Valid XML

- *Well-Formed XML* allows you to invent your own tags.
  - Similar to labels in semistructured data.
- *Valid XML* involves a DTD (*Document Type Definition*), a grammar for tags.

# Well-Formed XML

- Start the document with a *declaration*, surrounded by `<?xml ... ?>` .
- Normal declaration is:  

```
<?xml version = "1.0" standalone  
= "yes" ?>
```
- Balance of document is a *root tag* surrounding nested tags.

# Tags

- Tags, as in HTML, are normally matched pairs, as `<FOO> ... </FOO>` .
- Tags may be nested arbitrarily.
- XML tags are case sensitive.



# Example: Well-Formed XML

<?xml version = "1.0" standalone = "yes" ?>

<BARS>

<BAR><NAME>Joe's Bar</NAME>

A NAME  
subobject

<BEER><NAME>Bud</NAME>

<PRICE>2.50</PRICE></BEER>

<BEER><NAME>Miller</NAME>

<PRICE>3.00</PRICE></BEER>

A BEER  
subobject

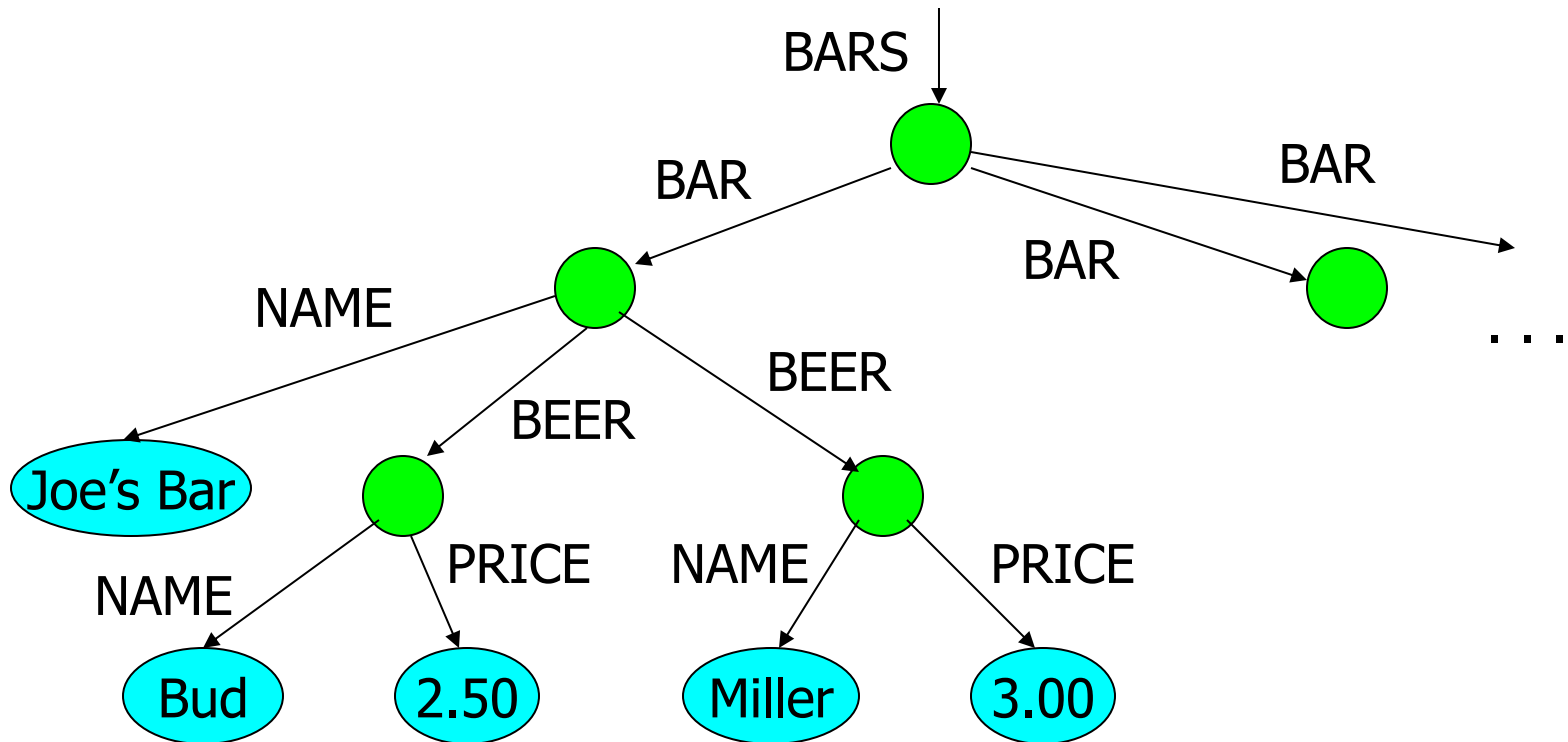
</BAR>

<BAR> ...

</BARS>

# Example

- Well-Formed XML with nested tags is exactly the same idea as trees of semistructured data.
- The <BARS> XML document is:



# XML Schema: DTD Structure

```
<!DOCTYPE <root tag> [  
    <!ELEMENT  <name> (<components>) >  
    ... more elements ...  

```

# XML Schema in XSD

```
<location>  
  <latitude>37.78333</latitude>  
  <longitude>122.4167</longitude>  
</location>
```

An XML schema for this element:

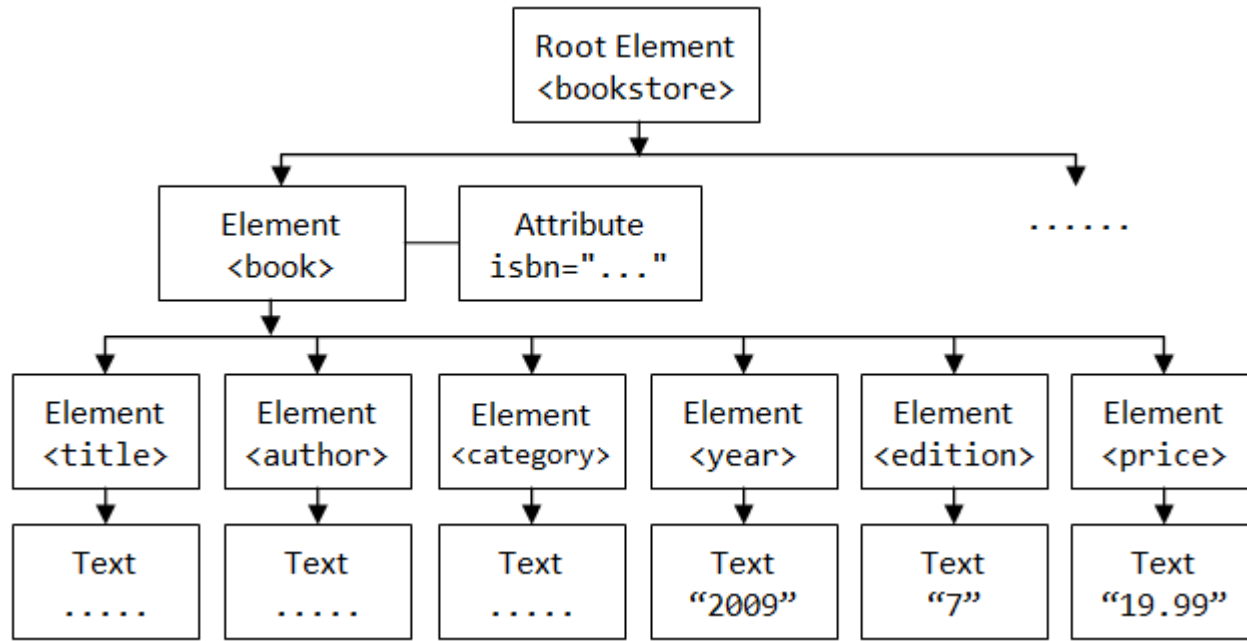
...

```
<xsd:schema xmlns:xsd="http://www.w3.org/2001/XMLSchema"  
  elementFormDefault="unqualified">  
  <xsd:complexType name="location">  
    <xsd:sequence>  
      <xsd:element name="latitude" type="xsd:decimal"/>  
      <xsd:element name="longitude" type="xsd:decimal"/>  
    </xsd:sequence>  
  </xsd:complexType name="location">
```

# XML and DOM

XML is a text format that encodes DOM (Document-Object Models) which is a data structure e.g. for Web pages.

The DOM is tree-structured:



# JSON format --> tree structure

```
{ "firstName": "John",  
  "lastName": "Smith",  
  "isAlive": true,  
  "age": 25,  
  "height_cm": 167.6,  
  "address": {  
    "streetAddress": "21 2nd Street",  
    "city": "New York",  
    "state": "NY",  
    "postalCode": "10021-3100"  
  }  
}
```

# Prerequisites for “Schemaless” DBs

- Need **external** and **internal** representations for all data types that will be used.
- **Internal:** a dynamically-typed, object-oriented language (like Java)
- **External:** an extensible data description language: JSON or XML
- **For Performance:** Fast SerDe (Serialization and DeSerialization) so internal data structures can be efficiently pushed or extracted from disk or network.

# Prerequisites for “Schemaless” DBs

- JSON includes named fields in a tree structure. Primitive types (e.g. string, number, boolean,...) are implicit.
- We can read JSON data (or XML) and automatically create internal representations for complex data.
- Using the **field names** and **object structure**, we can query these objects once loaded.



# XML vs. JSON

## XML:

- Separation between schema and data.
- Data can be represented and stored without schema (as strings).
- More verbose (but not true after compression or in DB).
- Standard Query/Transformation languages XSLT and Xquery.

## JSON:

- Types inferred inline. Schema rarely used but can be.
- Data without schema use type inference (string, int, float,...).
- More succinct in ASCII form.
- Transformation/ingestion rely on code (Java or Javascript).

# NoSQL Storage Systems



	Data Model
Cassandra	Columnfamily
CouchDB	Document
HBase	Columnfamily
MongoDB	Document
Neo4J	Graph
Redis	Collection
Riak	Document
Scalaris	Key/value
Tokyo Cabinet	Key/value
Voldemort	Key/value

# CouchDB Data Model (JSON)

- “With CouchDB, no schema is enforced, so new document types with new meaning can be safely added alongside the old.”
- A CouchDB document is an object that consists of named fields. Field values may be:
  - strings, numbers, dates,
  - ordered lists, associative maps

```
"Subject": "I like Plankton"  
"Author": "Rusty"  
"PostedDate": "5/23/2006"  
"Tags": ["plankton", "baseball", "decisions"]  
"Body": "I decided today that I don't like baseball. I like plankton."
```

# RDF Basics

- RDF is based on the idea of identifying resources using **Web identifiers** and describing resources in terms of simple **properties** and property **values**.
- To identify resources, RDF uses **Uniform Resource Identifiers (URIs)** and **URI references (URIrefs)**.
- **Definition:** A **resource** is anything that is identifiable by a URIref.

# Example

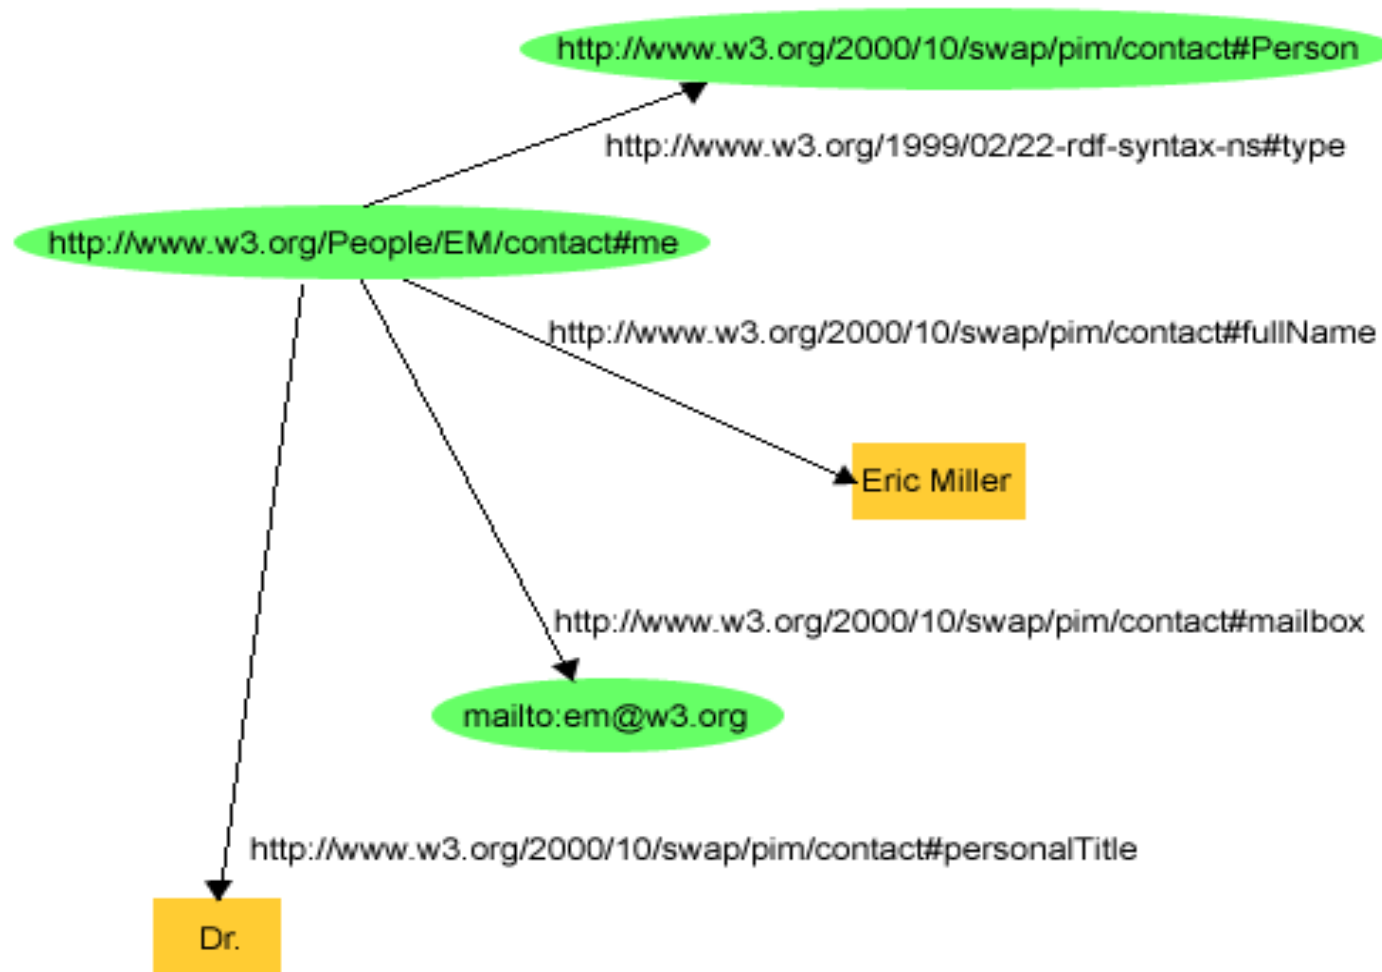
- Consider the following information:

“there is a Person identified by

`http://www.w3.org/People/EM/contact#me,`

whose name is Eric Miller, whose email address is `em@w3.org`, and whose title is Dr.”

# RDF Graph Example (cont'd)



# Basics (cont'd)

- Forget the long URIs for the moment!
- RDF is based on the idea that **the resources being described have properties which have values**, and that resources can be described by making **statements**, similar to the ones above, that specify those properties and values.
- **Terminology:**
  - The part that identifies the thing the statement is about is called the **subject**.
  - The part that identifies the property or characteristic of the subject that the statement specifies is called the **predicate**.
  - The part that identifies the value of that property is called the **object**.

# Example

`http://www.example.org/index.html` has  
a creator whose value is “John Smith”

- The **subject** is the URL

`http://www.example.org/index.html`

- The **predicate** is the word "creator"
- The **object** is the phrase “John Smith”



# RDF Triples → RDF Graphs (SPARQL)

- RDF statements can be written down using **triple notation**. In this notation, a statement is written as follows:

**subject predicate object .**

- **Example:**


```
<http://www.example.org/index.html>  
  <http://purl.org/dc/elements/1.1/creator>  
  <http://www.example.org/staffid/85740> .
```

```
<http://www.example.org/index.html>  
  <http://www.example.org/terms/creation-date> "August 16, 1999" .
```

```
<http://www.example.org/index.html>  
  <http://purl.org/dc/elements/1.1/language> "en" .
```

- **Note:** In this notation URIs are written out completely, in angle brackets.

# RDF Triples on the Web





# LOD: Linked RDF Triples on the Web

yago/wordnet:Artist109812338

yago/wordnet:Actor109765278

dfs:subclassOf

dfs:subclassOf

yago/wikicategory:ItalianComposer

imdb.com/name/nm0910607/

rdf:type

rdf:type

dbpedia.org/resource/Ennio\_Morricone

imdb.com/title/tt0361748/

op:composedMusicFor

dbpprop:citizenOf

dbpedia.org/resource/Rome

rdf.freebase.com/ns/en.rome

owl:sameAs

owl:sameAs

data.nytimes.com/51688803696189142301

geonames.org/3169070/roma

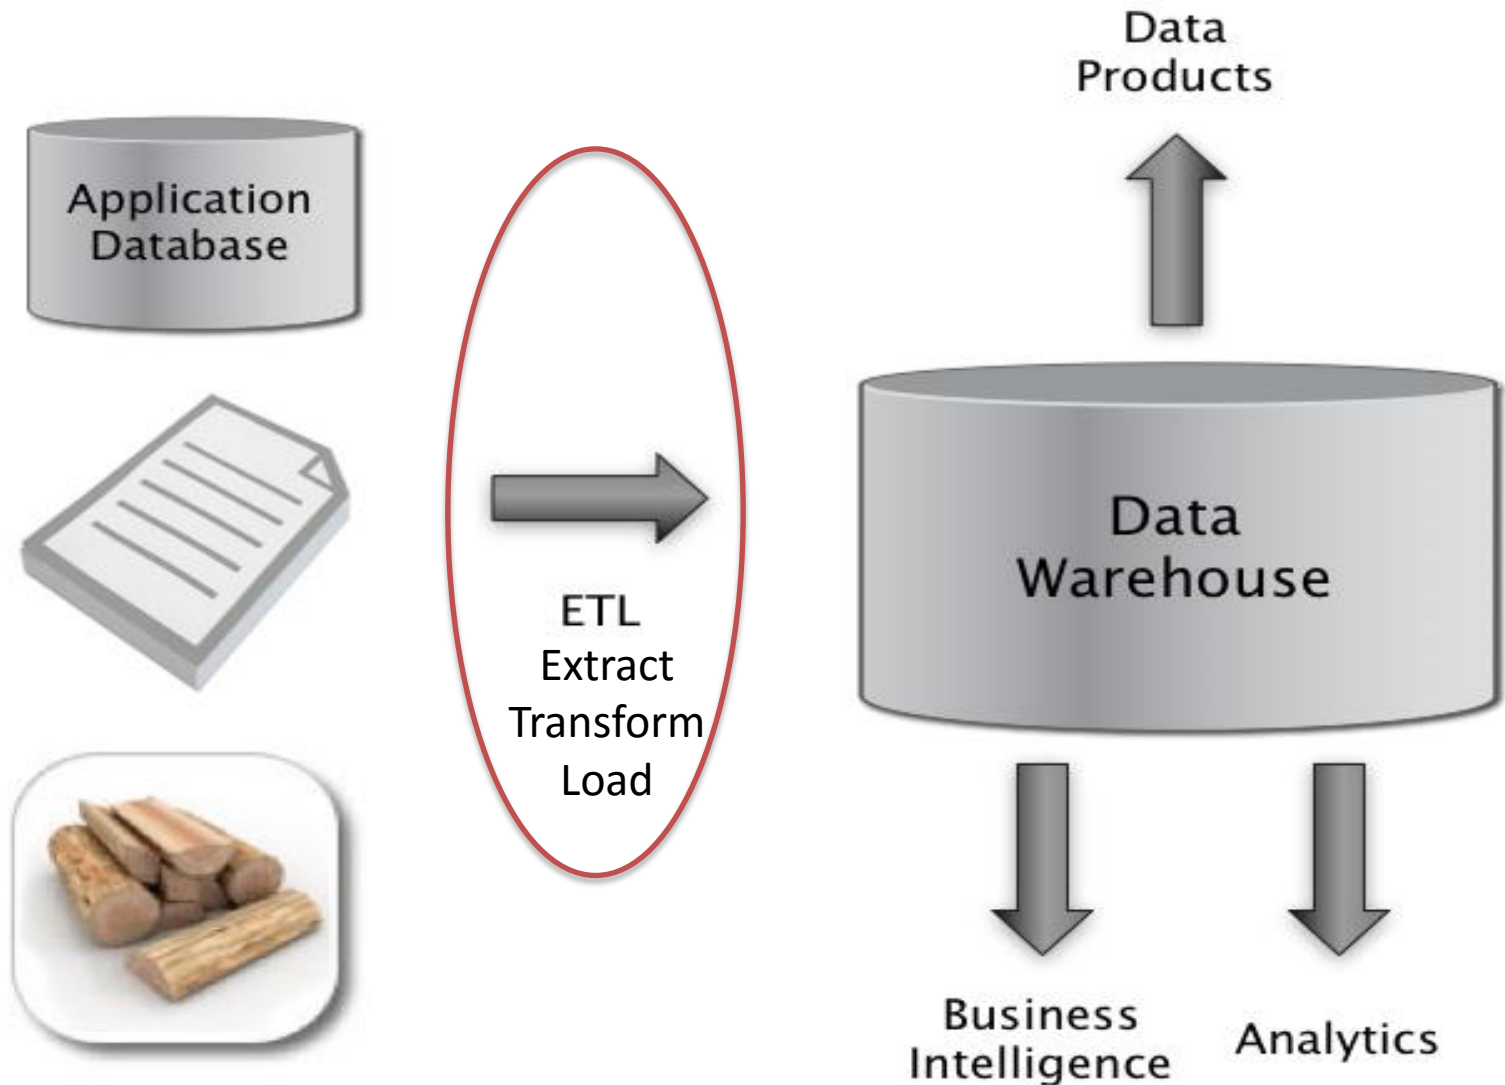
owl:sameAs

N 41° 54' 10" E 12° 29' 2"

# Outline

- Data Types and Sources
- Data Model
- Data Preparation
  - Big Picture

# The Big Picture



# What are your...

- Data Sources?
- ETL Process/Workflow and tools?
- Data Warehouse?
- Business Intelligence and Analytics?

# The Businessperson

- Data Sources
  - Web pages
  - Excel
- ETL
  - Copy and paste
- Data Warehouse
  - Excel
- Business Intelligence and Analytics
  - Excel functions
  - Excel charts

# The Programmer

- Data Sources
  - Web scraping, web services API
  - Excel spreadsheet exported as CSV
  - Database queries
- ETL
  - wget, curl, BeautifulSoup, lxml
- Data Warehouse
  - Flat files
- Business Intelligence and Analytics
  - Numpy, Matplotlib, R, Matlab



# The Enterprise

- Data Sources
  - Application databases
  - Intranet files
  - Application server log files
- ETL
  - Informatica, IBM DataStage, Ab Initio, Talend
- Data Warehouse
  - Teradata, Oracle, IBM DB2, Microsoft SQL Server
- Business Intelligence and Analytics
  - Business Objects, Cognos, Microstrategy
  - SAS, SPSS, R

# The Web Company

- Data Sources
  - Application databases
  - Logs from the services tier
  - Web crawl data
- ETL
  - Flume, Sqoop, Pig, Crunch, Oozie
- Data Warehouse
  - Hadoop/Hive, Spark/Shark
- Business Intelligence and Analytics
  - Custom dashboards: Argus, BirdBrain
  - R

# Data Preparation

- ETL
  - We need to **extract** data from the **source(s)**
  - We need to **load** data into the **sink**
  - We need to **transform** data at the source, sink, or in a **staging area**
  - Sources: file, database, event log, web site, HDFS...
  - Sinks: Python, R, SQLite, RDBMS, NoSQL store, files, HDFS...


# Data Preparation

- Workflow
  - The **pipeline** often consists of many steps
    - Using Unix command line pipes and filters
    - `head -50 wc_day6_1.log | cut -d ' ' -f 7 | sort | uniq -c | tail -10`
    - Simple data transformation using cut, sort, regex, sed, awk, etc.
  - Hands-on experience in Lab1
  - If the workflow is to be used more than once, it can be **scheduled**
    - Scheduling can be time-based or event-based
    - Use **publish-subscribe** to register interest (e.g. Twitter feeds)
  - Recording the execution of a workflow is known as capturing **lineage** or **provenance**

# HTML Tools - Parsing


- “Beautiful Soup”  
<http://www.crummy.com/software/BeautifulSoup/>  
a Python API for handling real HTML. DOM or SAX interfaces.
- “TagSoup”  
<http://ccil.org/~cowan/XML/tagsoup/>  
provides a Sax interface, i.e. a streaming parse, to Java applications. Can transform to a format you want using XSLT.
- Taggle, part of the Arabica toolset  
<http://www.jezuk.co.uk/cgi-bin/view/arabica/code>  
is a version of TagSoup written in C++. You may want to use this if you have a lot of data.

# Event-Driven Parsing: SAX

`<?xml version="1.0" encoding="UTF-8"?>`  Document Header

`<!-- bookstore.xml -->`  Comment

`<bookstore>`  Start-element “bookstore”

`<book ISBN="0123456001">`  Start-element “book”

`<title>Java For Dummies</title>`  Start-element “title”


`<author>Tan Ah Teck</author>` End-element “title”

`<category>Programming</category>`

`<year>2009</year>`

`<edition>7</edition>`

`<price>19.99</price>`

`</book>`  End-element “book”

# Event-Driven Parsing: SAX

A SAX parser finds all the **open-close-tag events** in an XML documents, and does **callbacks to user code**.

- User code can respond to only a subset of events corresponding to the tags it is interested in.
- User code can correctly compute aggregates from the data rather than create a record for each tag.
- User code must implement a state machine to keep track of “where it is” in the DOM tree.
- User code can implement flexible error recover strategies for ill-formed XML.
- Python XML parser: `xml.sax`, `lxml`

# What about JSON?

Most JSON parsers construct the “DOM” directly.

A few JSON (SAX-style) parsers:

- json, iJSON – Python
- Jackson – Java



# Summary

- Data Types and Sources
  - Schema on Read vs. Schema on Write
- Data Models
  - Set, HashTable, Tree, Graph..
- Data Preparation
- Unix data preparation lab 1 this Friday